

Does native FP4 compute help decode?

Stage A (measured, B300)

FP4 vs FP8 GEMM at M=128 (MLA decode)

Both HBM-bound: AI 100–176

FP4 $\leq$ 5.8% of 15 PF peak

FP8 $\leq$ 21.7% of 5 PF peak

KILL: FP4 compute never the limiter



Stage B (accuracy) /
Stage C (kernel)

not triggered

FP4/NVFP4 is a capacity + accuracy lever, not a decode-compute lever.