

Positioning: complement, not beat

We claim

- ✓ First open prediction-vs-measured decode roofline on sm_103
- ✓ MHA $\rightarrow$ GQA $\rightarrow$ MLA $\times$ FP16 $\rightarrow$ FP8 $\rightarrow$ NVFP4
- ✓ Per-CTA wall = work-starvation
- ✓ Two-layer model, predictions pre-registered before coding

We do NOT claim

- ✗ Faster than FlashInfer / FlashMLA / cuDNN
- ✗ A novel roofline method
- ✗ First to run on sm_103 (FA4 already deployed)
- ✗ A new SOTA kernel: v12 measured NVIDIA's own kernel