


The per-CTA wall was work-starvation

CUDA-core (v6–v11)

every lever → flat 0.75 TF



dead ends: TC, double-buffer, ILP, occupancy, bytes.

swap
engine
→

tcgen05 (v12)

2.4 → 1785 TF with work



36% FP8 peak / 46% HBM
at serving scale.